PRODUCT PITCH — ANTHROPIC

LakshX IDE

The agentic IDE that remembers — built for one developer today, an entire engineering team tomorrow.

PROTOTYPE — SHIPPED

ENTERPRISE AGENT BRAIN — IDEATION

lakshx.in MIT licensed VS Code fork

THE PROBLEM

Every AI coding agent today is a *brilliant amnesiac.*

`session-scoped`	Close the chat and the reasoning is gone. Tomorrow's session starts from zero, even on the same file.
`machine-scoped`	The agent only knows the one folder open on the one laptop it happens to be running on.
`unverified "done"`	The model **tells you** it finished. Nothing independently re-checks that claim against reality.

None of this is a model problem. It's an infrastructure problem.

OUR INSIGHT

The next unlock isn't a smarter model. It's *trustworthy autonomy.*

Verification, not vibes

"Done" is a re-run of a real, hash-frozen spec — a process exits 0, a test actually passes — not a sentence the model generated.

Persistent memory, not sessions

Every action the agent takes is traced and stored locally, always on — the record outlives the chat window.

Shared context, not one machine

The long-term bet: that memory shouldn't stop at one person's laptop when the work never did.

PRODUCT TODAY — SHIPPED

A VS Code fork with a real engineer inside it.

Royal Mode 2.0 SHIPPED

Plans, executes, and verifies its own work through a gated phase machine before it's allowed to call anything done.

Background subagents SHIPPED

Dispatches and monitors parallel work without blocking the main session.

Local trace store SHIPPED

Every agent action recorded locally and always on, no external observability service required.

Bring-your-own-model SHIPPED

Plug in any provider you already pay for — nothing locked to one vendor's API.

Structural search & secrets scan SHIPPED

AST-aware find/replace and SAST-lite security checks built into the editor.

Dependency & call graph SHIPPED

A guided tour of any unfamiliar codebase, generated automatically.

Voice mode SHIPPED

Offline, local speech-to-text push-to-talk dictation into the composer.

Remote access SHIPPED

Pair a phone over WiFi to watch — and steer — a running agent session.

● HOW IT EARNS TRUST

The agent can't just *say* it's done.

INTAKE → RECON + PLAN → EXECUTE → VERIFY → DONE

VERIFY fails → FIX (≤ 2 rounds) → back to EXECUTE

Still wrong → REWIND (≤ 2 re-entries) → back to PLAN

VERIFY re-runs a real, hash-frozen spec set before the plan was even executed — the model cannot quietly redefine "done" partway through a task. If it can't pass, it doesn't get to stop.

WHY NOT JUST CURSOR / COPILOT / WINDSURF?

They optimize the *typing.* We're building the *trust layer* underneath it.

	TYPICAL AI IDE	LAKSHX
Memory survives the session	— session only	✓ persistent, local
Independently verifies "done"	— self-reported	✓ hash-frozen spec re-run
Runs background multi-agent work	— single-threaded chat	✓ dispatch + monitor
Shares context across teammates	— isolated per machine	→ Agent Brain (roadmap)
Model-agnostic, local-first	— one vendor, cloud-only	✓ bring your own model

THE BIG SWING

Every agent today is isolated — to one employee, on one machine, inside one project folder.

Real engineering isn't done by one person. Why is the agent?

One team, five modules, five agents that have never met.

`auth dev` renames a token field — their agent knows, nobody else's does.

`api dev` ships an integration against the *old* field name — breaks silently in staging, three days later.

`qa engineer` files the bug. Her agent starts debugging from zero — it has no idea Tuesday's change even happened.

`new hire` joins mid-project and gets re-briefed by *humans* on decisions five agents already individually "knew."

One shared Agent Brain per project. Any teammate's agent can ask "has anyone touched this, and why?" and get a real, current answer — instead of a Slack archaeology dig.

THIS PATTERN IS EVERYWHERE

The five-person team is one example. It's not the only one.

Agencies with rotating developers

A contractor rolls off a client project after two months — today, everything their agent learned evaporates with them. Agent Brain persists at the project level, not the person level.

2 a.m. incident response

Whoever's on call starts cold. If this exact failure was already root-caused six weeks ago by a different agent, the Brain says so — and cuts time-to-resolution.

Regulated industries

Fintech and healthcare need a provable record of what changed, why, and how it was verified. The verification + trace infrastructure already produces exactly that.

Distributed, cross-timezone teams

A handoff today is a Slack message or a stale doc. With a shared Brain, the next timezone's agent already has full context on what happened eight hours ago.

THE BUSINESS CASE

Companies already pay for coordination. We're the first coding agent that reduces the bill instead of ignoring it.

already budgeted	Standups, sprint planning, code-review context-gathering, onboarding docs, tribal-knowledge risk on attrition — all real, existing line items on an engineering leader's budget.
compliance angle	Verification + trace infrastructure, already shipped, turns "an AI agent wrote this" from a liability into a searchable, provable audit trail.
who buys	This sells to a VP Eng or CTO for the **whole team** — not one developer's individual seat. Bigger contract, stickier, less churn.

NOT A PIVOT — AN EXTENSION

The hard infrastructure is already built. Agent Brain is what happens when we open it up.

SHIPPED, SOLO		TEAM-SCALE
Local trace store	→	Shared, project-level trace index
SHIPPED, SOLO		**TEAM-SCALE**
Background subagent registry	→	Cross-person task registry
SHIPPED, SOLO		**TEAM-SCALE**
Hash-frozen verification spec	→	Shared, team-wide "definition of done" per module

WHY WE CAN BUILD THIS

We didn't start with the vision slide. We started with the infrastructure it needs.

A shipped verification engine

Not a plan for one — running, tested, hash-frozen specs re-run against real process exit codes today.

Shipped multi-agent orchestration

Dispatching and monitoring parallel background work is the exact substrate a multiplayer Brain needs.

Shipped structured, persistent traces

The exact substrate shared memory needs — we just haven't opened it past one machine yet.

HOW THIS MAKES MONEY

Two tiers. Value scales with team size, not just seat count.

Individual LIVE TODAY

Seat-based. Bring your own model, run locally, keep every feature on this deck. The prototype you can install right now.

Team Agent Brain ROADMAP

Org-wide. Priced on team/project, not per head — the more people share one project's Brain, the more value it returns. Natural expansion revenue as teams grow.

WHERE WE ARE

Prototype. Ideation phase for the enterprise layer. Exactly where we say we are.

PROTOTYPE — SHIPPED → DESIGNING THE SHARED-CONTEXT PROTOCOL → TEAM PILOT → GENERAL AVAILABILITY

No overclaiming: the single-player product on slide 4 is real and running. Enterprise Agent Brain is the ideation-stage bet this raise is for.

BUILT WITH CLAUDE, BUILT FOR CLAUDE

This isn't a "we'll add Claude support" pitch. Claude is the default, top to bottom.

The product's own default model

A fresh LakshX install's out-of-the-box default provider is Anthropic, model Claude Sonnet 5 — not a configuration option buried in settings, the actual shipped default every new user starts from.

Built substantially with Claude Code itself

A real, verifiable share of this codebase's own commit history carries "Co-Authored-By: Claude Sonnet 5" — Royal Mode's verification engine, the background subagent registry, and this deck's own bug fixes were built in collaboration with Claude Code, not just marketed as AI-friendly.

The pitch: LakshX is a genuine showcase of Claude's agentic tool-use ceiling — verification-gated autonomy, background multi-agent orchestration, structured traces — running as the default experience, not a bolted-on integration.

WHAT WE'RE APPLYING FOR

API credits and rate limits to keep building — and a design partner as Royal Mode pushes on Claude's tool-use ceiling.

API credits + rate limits

To cover development traffic now, and the shared Agent Brain service's usage as it moves from ideation to a real pilot — both run on Claude by default, so this is directly proportional to Anthropic's own usage, not a side request.

Early access + a real feedback loop

Royal Mode's verification-gated autonomy and background subagent orchestration push hard on tool-use and long-horizon agentic behavior — exactly the surface area new Claude capabilities land on first. Happy to be a design partner / early-access tester and report back with real, shipped-product data, not a lab benchmark.

Raise amount, equity, and cohort details: fill in before submitting — deliberately left open here.

WHERE THIS GOES

From autocomplete. To an autonomous agent. To an engineering team's *shared brain.*