



LakshX_{IDE}

*The agentic IDE that remembers — built
for one developer today, an entire
engineering team tomorrow.*

• PROTOTYPE — SHIPPED

• ENTERPRISE AGENT BRAIN — IDEATION

● THE PROBLEM

Every AI coding agent today is a *brilliant amnesiac*.

- `session-scoped` Close the chat and the reasoning is gone. Tomorrow's session starts from zero, even on the same file.
- `machine-scoped` The agent only knows the one folder open on the one laptop it happens to be running on.
- `unverified "done"` The model **tells you** it finished. Nothing independently re-checks that claim against reality.

None of this is a model problem. It's an infrastructure problem.

● OUR INSIGHT

The next unlock isn't a smarter model. It's *trustworthy autonomy*.

Verification, not vibes

"Done" is a re-run of a real, hash-frozen spec — a process exits 0, a test actually passes — not a sentence the model generated.

Persistent memory, not sessions

Every action the agent takes is traced and stored locally, always on — the record outlives the chat window.

Shared context, not one machine

The long-term bet: that memory shouldn't stop at one person's laptop when the work never did.

A VS Code fork with a real engineer inside it.

Royal Mode 2.0

SHIPPED

Plans, executes, and verifies its own work through a gated phase machine before it's allowed to call anything done.

Background subagents

SHIPPED

Dispatches and monitors parallel work without blocking the main session.

Local trace store

SHIPPED

Every agent action recorded locally and always on, no external observability service required.

Bring-your-own-model

SHIPPED

Plug in any provider you already pay for — nothing locked to one vendor's API.

Structural search & secrets scan

SHIPPED

AST-aware find/replace and SAST-lite security checks built into the editor.

Dependency & call graph

SHIPPED

A guided tour of any unfamiliar codebase, generated automatically.

Voice mode

SHIPPED

Offline, local speech-to-text push-to-talk dictation into the composer.

Remote access

SHIPPED

Pair a phone over WiFi to watch — and steer — a running agent session.

● HOW IT EARNS TRUST

The agent can't just *say* it's done.



VERIFY fails → FIX (≤ 2 rounds) → back to EXECUTE

Still wrong → REWIND (≤ 2 re-entries) → back to PLAN

VERIFY re-runs a real, hash-frozen spec set before the plan was even executed — the model cannot quietly redefine "done" partway through a task. If it can't pass, it doesn't get to stop.

● WHY NOT JUST CURSOR / COPILOT / WINDSURF?

They optimize the *typing*. We're building the *trust layer* underneath it.

	TYPICAL AI IDE	LAKSHX
Memory survives the session	– session only	✓ persistent, local
Independently verifies "done"	– self-reported	✓ hash-frozen spec re-run
Runs background multi-agent work	– single-threaded chat	✓ dispatch + monitor
Shares context across teammates	– isolated per machine	→ Agent Brain (roadmap)
Model-agnostic, local-first	– one vendor, cloud-only	✓ bring your own model

● THE BIG SWING

Every agent today
is isolated — to one
employee, on one
machine, inside
one project folder.

Real engineering isn't
done by one person.
Why is the agent?

One team, five modules, five agents that have never met.

auth dev	renames a token field — their agent knows, nobody else's does.
api dev	ships an integration against the <i>old</i> field name — breaks silently in staging, three days later.
qa engineer	files the bug. Her agent starts debugging from zero — it has no idea Tuesday's change even happened.
new hire	joins mid-project and gets re-briefed by <i>humans</i> on decisions five agents already individually "knew."



One shared Agent Brain per project. Any teammate's agent can ask "has anyone touched this, and why?" and get a real, current answer — instead of a Slack archaeology dig.

● THIS PATTERN IS EVERYWHERE

The five-person team is one example. It's not the only one.

Agencies with rotating developers

A contractor rolls off a client project after two months — today, everything their agent learned evaporates with them. Agent Brain persists at the project level, not the person level.

2 a.m. incident response

Whoever's on call starts cold. If this exact failure was already root-caused six weeks ago by a different agent, the Brain says so — and cuts time-to-resolution.

Regulated industries

Fintech and healthcare need a provable record of what changed, why, and how it was verified. The verification + trace infrastructure already produces exactly that.

Distributed, cross-timezone teams

A handoff today is a Slack message or a stale doc. With a shared Brain, the next timezone's agent already has full context on what happened eight hours ago.

● THE BUSINESS CASE

Companies already pay for coordination.
We're the first coding agent that
reduces the bill instead of ignoring it.

already
budgeted

Standups, sprint planning, code-review context-gathering, onboarding docs, tribal-knowledge risk on attrition — all real, existing line items on an engineering leader's budget.

compliance
angle

Verification + trace infrastructure, already shipped, turns "an AI agent wrote this" from a liability into a searchable, provable audit trail.

who buys

This sells to a VP Eng or CTO for the **whole team** — not one developer's individual seat. Bigger contract, stickier, less churn.

● NOT A PIVOT – AN EXTENSION

The hard infrastructure is already built. Agent Brain is what happens when we open it up.

SHIPPED, SOLO

Local trace store



TEAM-SCALE

Shared, project-level trace index

SHIPPED, SOLO

Background subagent registry



TEAM-SCALE

Cross-person task registry

SHIPPED, SOLO

Hash-frozen verification spec



TEAM-SCALE

Shared, team-wide "definition of done" per module

● WHY WE CAN BUILD THIS

We didn't start with the vision slide. We started with the infrastructure it needs.

A shipped verification engine

Not a plan for one — running, tested, hash-frozen specs re-run against real process exit codes today.

Shipped multi-agent orchestration

Dispatching and monitoring parallel background work is the exact substrate a multiplayer Brain needs.

Shipped structured, persistent traces

The exact substrate shared memory needs — we just haven't opened it past one machine yet.

● HOW THIS MAKES MONEY

Two tiers. Value scales with team size, not just seat count.

Individual

LIVE TODAY

Seat-based. Bring your own model, run locally, keep every feature on this deck. The prototype you can install right now.


Team Agent Brain

ROADMAP

Org-wide. Priced on team/project, not per head — the more people share one project's Brain, the more value it returns. Natural expansion revenue as teams grow.

● WHERE WE ARE

Prototype. Ideation phase for the enterprise layer. Exactly where we say we are.



No overclaiming: the single-player product on slide 4 is real and running. Enterprise Agent Brain is the ideation-stage bet this raise is for.

● WHAT WE'RE APPLYING FOR

Fuel for the single-player product's growth, and the infrastructure the multiplayer bet needs.

YC

Speed and rigor on GTM — turning five design-partner teams into a repeatable enterprise motion.

MICROSOFT FOR STARTUPS

Azure credits to host the shared Agent Brain service, plus a natural distribution channel through VS Code's own ecosystem.

AWS ACTIVATE

Compute credits for the trace/index infrastructure a team-scale Brain runs on, at the scale a pilot cohort needs.

Raise amount, equity, and cohort details: fill in before submitting — deliberately left open here.

● WHERE THIS GOES

From autocomplete.
To an autonomous
agent. To an
engineering
team's *shared brain*.